

Healthcare

October 15, 2022

```
[1]: #Import Lybraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: data=pd.read_csv('health care diabetes.csv')
```

```
[3]: data.head()
```

```
[3]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0             6     148             72             35         0  33.6
1             1      85             66             29         0  26.6
2             8     183             64              0         0  23.3
3             1      89             66             23        94  28.1
4             0     137             40             35       168  43.1

      DiabetesPedigreeFunction  Age  Outcome
0                0.627     50         1
1                0.351     31         0
2                0.672     32         1
3                0.167     21         0
4                2.288     33         1
```

```
[4]: data.isnull().any()
```

```
[4]: Pregnancies      False
Glucose              False
BloodPressure        False
SkinThickness        False
Insulin              False
BMI                  False
DiabetesPedigreeFunction  False
Age                  False
Outcome              False
dtype: bool
```

```
[5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null   int64
1   Glucose                768 non-null   int64
2   BloodPressure          768 non-null   int64
3   SkinThickness          768 non-null   int64
4   Insulin                768 non-null   int64
5   BMI                   768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                   768 non-null   int64
8   Outcome                768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```


0.1 Using Histograms

```
[6]: data['Glucose'].value_counts().head(7)
```

```
[6]: 100    17
     99    17
     129   14
     125   14
     111   14
     106   14
     95    13
     Name: Glucose, dtype: int64
```

```
[7]: plt.hist(data['Glucose'])
```

```
[7]: (array([ 5.,  0.,  4., 32., 156., 211., 163., 95., 56., 46.]),
      array([ 0., 19.9, 39.8, 59.7, 79.6, 99.5, 119.4, 139.3, 159.2,
              179.1, 199. ]),
      <BarContainer object of 10 artists>)
```




```
[8]: data['BloodPressure'].value_counts().head(7)
```

```
[8]: 70    57
      74    52
      68    45
      78    45
      72    44
      64    43
      80    40
      Name: BloodPressure, dtype: int64
```

```
[9]: plt.hist(data['BloodPressure'])
```

```
[9]: (array([ 35.,  1.,  2., 13., 107., 261., 243.,  87., 14.,  5.]),
      array([ 0. , 12.2, 24.4, 36.6, 48.8, 61. , 73.2, 85.4, 97.6,
              109.8, 122. ]),
      <BarContainer object of 10 artists>)
```




```
[10]: data['SkinThickness'].value_counts().head(7)
```

```
[10]: 0      227
      32      31
      30      27
      27      23
      23      22
      33      20
      18      20
      Name: SkinThickness, dtype: int64
```

```
[11]: plt.hist(data['SkinThickness'])
```

```
[11]: (array([231., 107., 165., 175., 78., 9., 2., 0., 0., 1.]),
      array([ 0. , 9.9, 19.8, 29.7, 39.6, 49.5, 59.4, 69.3, 79.2, 89.1, 99. ]),
      <BarContainer object of 10 artists>)
```




```
[12]: data['Insulin'].value_counts().head(7)
```

```
[12]: 0      374
      105     11
      140      9
      130      9
      120      8
      100      7
       94      7
      Name: Insulin, dtype: int64
```

```
[13]: plt.hist(data['Insulin'])
```

```
[13]: (array([487., 155., 70., 30., 8., 9., 5., 1., 2., 1.]),
      array([ 0. , 84.6, 169.2, 253.8, 338.4, 423. , 507.6, 592.2, 676.8,
              761.4, 846. ]),
      <BarContainer object of 10 artists>)
```




```
[14]: data['BMI'].value_counts().head(7)
```

```
[14]: 32.0    13
      31.6    12
      31.2    12
      0.0    11
      33.3    10
      32.4    10
      32.8     9
      Name: BMI, dtype: int64
```

```
[15]: plt.hist(data['BMI'])
```

```
[15]: (array([ 11.,   0.,  15., 156., 268., 224.,  78.,  12.,   3.,   1.]),
      array([ 0. ,   6.71, 13.42, 20.13, 26.84, 33.55, 40.26, 46.97, 53.68,
              60.39, 67.1 ]),
      <BarContainer object of 10 artists>)
```



```
[16]: data.describe().transpose()
```

```
[16]:
```

	count	mean	std	min	25%	\
Pregnancies	768.0	3.845052	3.369578	0.000	1.00000	
Glucose	768.0	120.894531	31.972618	0.000	99.00000	
BloodPressure	768.0	69.105469	19.355807	0.000	62.00000	
SkinThickness	768.0	20.536458	15.952218	0.000	0.00000	
Insulin	768.0	79.799479	115.244002	0.000	0.00000	
BMI	768.0	31.992578	7.884160	0.000	27.30000	
DiabetesPedigreeFunction	768.0	0.471876	0.331329	0.078	0.24375	
Age	768.0	33.240885	11.760232	21.000	24.00000	
Outcome	768.0	0.348958	0.476951	0.000	0.00000	

	50%	75%	max
Pregnancies	3.0000	6.00000	17.00
Glucose	117.0000	140.25000	199.00
BloodPressure	72.0000	80.00000	122.00
SkinThickness	23.0000	32.00000	99.00
Insulin	30.5000	127.25000	846.00
BMI	32.0000	36.60000	67.10
DiabetesPedigreeFunction	0.3725	0.62625	2.42
Age	29.0000	41.00000	81.00
Outcome	0.0000	1.00000	1.00

```
[17]: Positive=data[data['Outcome']==1]
      Positive.head(5)
```


```
[17]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI   \
0           6        148             72              35         0  33.6
2           8        183             64              0         0  23.3
4           0        137             40              35        168  43.1
6           3         78             50              32         88  31.0
8           2        197             70              45        543  30.5

      DiabetesPedigreeFunction  Age  Outcome
0                0.627      50         1
2                0.672      32         1
4                2.288      33         1
6                0.248      26         1
8                0.158      53         1
```

0.2 Using histograms treat the missing values

```
[18]: plt.hist(Positive['Glucose'],histtype='stepfilled',bins=20)
```

```
[18]: (array([ 2.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  4.,  9., 28., 26., 36.,
        27., 29., 22., 24., 21., 25., 14.]),
      array([ 0. ,  9.95, 19.9 , 29.85, 39.8 , 49.75, 59.7 , 69.65,
        79.6 , 89.55, 99.5 , 109.45, 119.4 , 129.35, 139.3 , 149.25,
        159.2 , 169.15, 179.1 , 189.05, 199.  ]),
      [<matplotlib.patches.Polygon at 0x7f2aad9fa50>])
```





```
[19]: Positive['Glucose'].value_counts().head(7)
```

```
[19]: 125    7
      158    6
      128    6
      115    6
      129    6
      146    5
      162    5
      Name: Glucose, dtype: int64
```

```
[20]: plt.hist(Positive['BloodPressure'],histtype='stepfilled',bins=20)
```

```
[20]: (array([16.,  0.,  0.,  0.,  0.,  1.,  0.,  1.,  6.,  6., 19., 37., 56.,
          36., 41., 31.,  7.,  4.,  4.,  3.]),
      array([ 0.,  5.7, 11.4, 17.1, 22.8, 28.5, 34.2, 39.9, 45.6,
          51.3, 57., 62.7, 68.4, 74.1, 79.8, 85.5, 91.2, 96.9,
          102.6, 108.3, 114. ]),
      [<matplotlib.patches.Polygon at 0x7f2aadc2ead0>])
```




```
[21]: Positive['BloodPressure'].value_counts().head(7)
```

```
[21]: 70    23
      76    18
      78    17
      74    17
```

```
72    16
0      16
82     13
Name: BloodPressure, dtype: int64
```

```
[22]: plt.hist(Positive['SkinThickness'],histtype='stepfilled',bins=20)
```

```
[22]: (array([88.,  1.,  4., 10., 18., 30., 41., 34., 23., 15.,  1.,  1.,  1.,
          0.,  0.,  0.,  0.,  0.,  0.,  1.]),
      array([ 0.   ,  4.95,  9.9 , 14.85, 19.8 , 24.75, 29.7 , 34.65, 39.6 ,
          44.55, 49.5 , 54.45, 59.4 , 64.35, 69.3 , 74.25, 79.2 , 84.15,
          89.1 , 94.05, 99.   ]),
      [<matplotlib.patches.Polygon at 0x7f2aadb82d0>])
```




```
[23]: Positive['SkinThickness'].value_counts().head(7)
```

```
[23]: 0      88
      32     14
      33      9
      30      9
      39      8
      35      8
      36      8
Name: SkinThickness, dtype: int64
```

```
[24]: plt.hist(Positive['Insulin'],histtype='stepfilled',bins=20)
```

```
[24]: (array([141.,  6., 23., 33., 24., 12.,  7.,  7.,  2.,  1.,  1.,
           5.,  3.,  1.,  1.,  0.,  0.,  0.,  0.,  1.]),
       array([ 0. , 42.3, 84.6, 126.9, 169.2, 211.5, 253.8, 296.1, 338.4,
           380.7, 423. , 465.3, 507.6, 549.9, 592.2, 634.5, 676.8, 719.1,
           761.4, 803.7, 846. ]),
       [<matplotlib.patches.Polygon at 0x7f2aadb9bf50>])
```




```
[25]: Positive['Insulin'].value_counts().head(7)
```

```
[25]: 0      138
      130      6
      180      4
      156      3
      175      3
      194      2
      125      2
      Name: Insulin, dtype: int64
```

```
[26]: plt.hist(Positive['BMI'],histtype='stepfilled',bins=20)
```

```
[26]: (array([ 2.,  0.,  0.,  0.,  0.,  0.,  3., 13., 38., 61., 61., 36., 27.,
           14.,  7.,  3.,  1.,  1.,  0.,  1.]),
       array([ 0.   ,  3.355,  6.71 , 10.065, 13.42 , 16.775, 20.13 , 23.485,
           26.84 , 30.195, 33.55 , 36.905, 40.26 , 43.615, 46.97 , 50.325,
           53.68 , 57.035, 60.39 , 63.745, 67.1  ]),
       [<matplotlib.patches.Polygon at 0x7f2aadaa80d0>])
```



```
[27]: Positive['BMI'].value_counts().head(7)
```

```
[27]: 32.9      8
      31.6      7
      33.3      6
      30.5      5
      32.0      5
      31.2      5
      32.4      4
      Name: BMI, dtype: int64
```

```
[28]: Positive.describe().transpose()
```

```
[28]:
```

	count	mean	std	min	25%	\
Pregnancies	268.0	4.865672	3.741239	0.000	1.7500	
Glucose	268.0	141.257463	31.939622	0.000	119.0000	
BloodPressure	268.0	70.824627	21.491812	0.000	66.0000	
SkinThickness	268.0	22.164179	17.679711	0.000	0.0000	
Insulin	268.0	100.335821	138.689125	0.000	0.0000	
BMI	268.0	35.142537	7.262967	0.000	30.8000	
DiabetesPedigreeFunction	268.0	0.550500	0.372354	0.088	0.2625	
Age	268.0	37.067164	10.968254	21.000	28.0000	
Outcome	268.0	1.000000	0.000000	1.000	1.0000	


	50%	75%	max
Pregnancies	4.000	8.000	17.00

Glucose	140.000	167.000	199.00
BloodPressure	74.000	82.000	114.00
SkinThickness	27.000	36.000	99.00
Insulin	0.000	167.250	846.00
BMI	34.250	38.775	67.10
DiabetesPedigreeFunction	0.449	0.728	2.42
Age	36.000	44.000	70.00
Outcome	1.000	1.000	1.00


1 Scatter Plot

```
[29]: BloodPressure= Positive['BloodPressure']
      Glucose=       Positive['Glucose']
      SkinThickness= Positive['SkinThickness']
      Insulin=       Positive['Insulin']
      BMI=           Positive['BMI']
```


```
[30]: plt.scatter(BloodPressure,Glucose,color=['blue'])
      plt.xlabel('BloodPressure')
      plt.ylabel('Glucose')
      plt.title('Glucose & BloodPressure')
      plt.show()
```




```
[31]: g = sns.scatterplot(x='Glucose', y='BloodPressure',  
                        hue='Outcome',  
                        data=data);
```



```
[32]: B = sns.scatterplot(x='BMI', y='Insulin',  
                        hue='Outcome',  
                        data=data);
```



```
[33]: g = sns.scatterplot(x='SkinThickness', y='Insulin',  
                        hue='Outcome',  
                        data=data);
```



1.1 Correlation analysis

```
[34]: data.corr()
```

```
[34]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	\
Pregnancies	1.000000	0.129459	0.141282	-0.081672	
Glucose	0.129459	1.000000	0.152590	0.057328	
BloodPressure	0.141282	0.152590	1.000000	0.207371	
SkinThickness	-0.081672	0.057328	0.207371	1.000000	
Insulin	-0.073535	0.331357	0.088933	0.436783	
BMI	0.017683	0.221071	0.281805	0.392573	
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928	
Age	0.544341	0.263514	0.239528	-0.113970	
Outcome	0.221898	0.466581	0.065068	0.074752	


	Insulin	BMI	DiabetesPedigreeFunction	\
Pregnancies	-0.073535	0.017683	-0.033523	
Glucose	0.331357	0.221071	0.137337	
BloodPressure	0.088933	0.281805	0.041265	
SkinThickness	0.436783	0.392573	0.183928	
Insulin	1.000000	0.197859	0.185071	
BMI	0.197859	1.000000	0.140647	
DiabetesPedigreeFunction	0.185071	0.140647	1.000000	
Age	-0.042163	0.036242	0.033561	
Outcome	0.130548	0.292695	0.173844	

	Age	Outcome
Pregnancies	0.544341	0.221898
Glucose	0.263514	0.466581
BloodPressure	0.239528	0.065068
SkinThickness	-0.113970	0.074752
Insulin	-0.042163	0.130548
BMI	0.036242	0.292695
DiabetesPedigreeFunction	0.033561	0.173844
Age	1.000000	0.238356
Outcome	0.238356	1.000000

2 Heatmaps


```
[35]: sns.heatmap(data.corr())
```

```
[35]: <AxesSubplot:>
```

```
[36]: plt.subplots(figsize=(8,8))
      sns.heatmap(data.corr(),annot=True,cmap='viridis')
```


[36]: <AxesSubplot:>



```
[37]: plt.subplots(figsize=(8,8))
      sns.heatmap(data.corr(),annot=True)

      # Given Correlation Values
```

```
[37]: <AxesSubplot:>
```



3 Week 2 / Data Modeling

```
[38]: # Logistic Regreesion & Model Building
```

```
data.head()
```

```
[38]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	

3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[39]: #Train Test Split
features=data.iloc[:,[0,1,2,3,4,5,6,7,]].values
label= data.iloc[:,8].values
```

```
[40]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(features,
                                                label,
                                                test_size=0.2,
                                                random_state=10)
```

```
[41]: from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(X_train,y_train)
```

```
[41]: LogisticRegression()
```

```
[42]: print(model.score(X_train,y_train))
print(model.score(X_test,y_test))
```

```
0.7719869706840391
0.7662337662337663
```

```
[43]: from sklearn.metrics import confusion_matrix
cm = confusion_matrix(label,model.predict(features))
cm
```

```
[43]: array([[446,  54],
          [122, 146]])
```

```
[44]: from sklearn.metrics import classification_report
print(classification_report(label,model.predict(features)))
```

	precision	recall	f1-score	support
0	0.79	0.89	0.84	500
1	0.73	0.54	0.62	268
accuracy			0.77	768


macro avg	0.76	0.72	0.73	768
weighted avg	0.77	0.77	0.76	768

```
[45]: #Preparing ROC Curve(Receiver Operating Characteristics Curve)
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

#Predict Probabilities
probs = model.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
auc = roc_auc_score(label, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
fpr, tpr, thresholds = roc_curve(label, probs)
# plot no skill
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(fpr, tpr, marker='.')
```

AUC: 0.837

[45]: [<matplotlib.lines.Line2D at 0x7f2aa65b8d50>]



```
[46]: #Applying Decission Tree Classifier  
from sklearn.tree import DecisionTreeClassifier  
model3 = DecisionTreeClassifier(max_depth=5)  
model3.fit(X_train,y_train)
```

```
[46]: DecisionTreeClassifier(max_depth=5)
```

```
[47]: model3.score(X_train,y_train)
```

```
[47]: 0.8289902280130294
```

```
[48]: model3.score(X_test,y_test)
```

```
[48]: 0.7597402597402597
```

```
[49]: #Applying Random Forest  
from sklearn.ensemble import RandomForestClassifier  
model4= RandomForestClassifier(n_estimators=1)  
model4.fit(X_train,y_train)
```

```
[49]: RandomForestClassifier(n_estimators=1)
```

```
[50]: model4.score(X_train,y_train)
```

```
[50]: 0.8990228013029316
```

```
[51]: model4.score(X_test,y_test)
```

```
[51]: 0.7077922077922078
```

```
[52]: #Support Vector Classifier  
  
from sklearn.svm import SVC  
model5 = SVC(kernel='rbf',  
              gamma='auto')  
model5.fit(X_train,y_train)
```

```
[52]: SVC(gamma='auto')
```

```
[53]: model5.score(X_test,y_test),model5.score(X_train,y_train)
```

```
[53]: (0.6168831168831169, 1.0)
```

```
[54]: #Applying K-NN  
from sklearn.neighbors import KNeighborsClassifier  
model2 = KNeighborsClassifier(n_neighbors=7,  
                             metric='minkowski',
```

```

p = 2)
model2.fit(X_train,y_train)

```

```
[54]: KNeighborsClassifier(n_neighbors=7)
```

```

[55]: #Preparing ROC Curve (Receiver Operating Characteristics Curve)
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

# predict probabilities
probs = model2.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
auc = roc_auc_score(label, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
fpr, tpr, thresholds = roc_curve(label, probs)
print("True Positive Rate - {}, False Positive Rate - {} Thresholds - {}".
      ↪format(tpr,fpr,thresholds))
# plot no skill
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(fpr, tpr, marker='.')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")

```


AUC: 0.836

```

True Positive Rate - [0.      0.06716418 0.23880597 0.44776119 0.60074627
0.75373134
0.88059701 0.98507463 1.      ], False Positive Rate - [0.      0.      0.02
0.056 0.12  0.248 0.428 0.668 1.      ] Thresholds - [2.      1.
0.85714286 0.71428571 0.57142857 0.42857143
0.28571429 0.14285714 0.      ]

```

```
[55]: Text(0, 0.5, 'True Positive Rate')
```




```
[56]: #Precision Recall Curve for Logistic Regression

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')

```


f1=0.624 auc=0.726 ap=0.727

[56]: [<matplotlib.lines.Line2D at 0x7f2aa591de50>]




```
[57]: #Precision Recall Curve for KNN Algorithms

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model2.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model2.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
```

```
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')
```

f1=0.658 auc=0.752 ap=0.709

[57]: [<matplotlib.lines.Line2D at 0x7f2aa58a9810>]



[58]: *#Precision Recall Curve for Decission Tree Classifier*

```
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score

# predict probabilities
probs = model3.predict_proba(features)

# keep probabilities for the positive outcome only
probs = probs[:, 1]

# predict class values
#yhat = model3.predict(features)

# calculate precision-recall curve
#precision, recall, thresholds = precision_recall_curve(label, probs)
```

```

# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)

# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))


# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')

# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')

```

f1=0.658 auc=0.752 ap=0.761

[58]: [<matplotlib.lines.Line2D at 0x7f2aa582bf90>]



[59]: *#Precision Recall Curve for Random Forest*

```

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score

```


```

# predict probabilities
probs = model4.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model4.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')

```

f1=0.803 auc=0.836 ap=0.710

[59]: [<matplotlib.lines.Line2D at 0x7f2aa57b2b50>]



3.0.1 Thank You

[]: